

shared_ptr <T>

Sharing ownership of an object through a pointer

deallocated < last remaining shared_ptr is destroyed (delete)
< operator= or reset()

make_shared <T>

creates a shared pointer that manages a new object

i.e: shared_ptr <int> foo = make_shared <int> (10);

tcdrain()

used to block until all output written to the object is transmitted

shared_lock()

Balance readers and writers. No one will have higher priority here.

lock_guard()

you don't need to call lock() and unlock() manually

- it is unsafe for a member function to return a pointer or reference

deadlock:

a situation where two or more threads are blocked forever, waiting for each other. (Nothing gonna happen)

race condition:

occurs when multiple concurrently access a shared data

Rather than using an instance of std::mutex for the synchronization, you use an instance of boost::shared_mutex. For the update operations, std::lock_guard<boost::shared_mutex> and std::unique_lock<boost::shared_mutex> can be used for the locking, in place of the corresponding std::mutex specializations. These ensure exclusive access, just as with std::mutex. Those threads that don't need to update the data structure can instead use boost::shared_lock<boost::shared_mutex>

Collaboration Graph Notation (CGN)

```
/*
 * Fraction.cpp
 *
 * Multiple threads accessing a simple fraction object
 *
 * 2019 C. Scratchley created to work with C++11/14
 *
 * Copyright 2019, School of Engineering Science, SFU, Canada
 */
```

Oct 7th recording

copy mutex

delete.

```
#include <iostream>
#include <signal.h>
#include <memory>
#include <sched.h> // sched_yield()
#include <unistd.h> /* required for sleep() */
#include <stdlib.h> /* abs() ? */
#include <thread>
#include <mutex> // std::mutex.h // mutual exclusion
```

Chapter 3

using namespace std;

help to fix the code

where should we put a mutex?

```
mutex consoleMutex;
#define COUT (lock_guard<mutex>(consoleMutex), std::cout)
```

class fraction

```
{ private: mutex myMutex; } // good
public:
    unsigned num;
    unsigned denom;
```

either public or private.

use lock mutex to have only 1 thread running.

// bad, 2 threads may be working on 2 different vars, and mutex will prevent it.

```
fraction(unsigned numerator, unsigned denominator) {
```

num = numerator;
 denom = denominator;

};

```
void report(char id) {
    lock_guard<mutex> myMutex{myMutex}; // copy
```

lock - guard guard{ myMutex }

+ protect two lines data below

unwanted

b/w threads

```
    cout << id << ": Numerator is " << numC << " Denominator is " << denomC
        << ". Floating representation is " << 1.0 * numC / denomC
        << ". Integer representation is " << numC / denomC
        << endl;
```

void invert()

mutex myMutex; // bad because each caller gets own mutex

```
}; // lock_guard guard(myMutex)
```

```
void invert() { //  $\frac{2}{3} \rightarrow \frac{3}{2}$ 
    lock_guard<mutex> myMutex{myMutex};
    myMutex.lock();
```

A int diff = num - denom; // subtract num count from denom count

A if (diff >= 0) {

```
// A B sched_yield(); // may cause crash. the program is changing data,
    num -= diff; // remove from num
    denom += diff; // and... insert into denom
};
```

jump to other funcs, can cause wrong data.

```

else {
    unsigned adiff = abs(diff);
    denom -= adiff; // remove from denom
    num += adiff; // and... insert into num
};

// unlock the mutex myMutex.unlock();
};

//fraction fraction1(2, 3);
fraction fraction1(3, 6); & global object

void threadFunc(char id) {
    // while (true) {
        for (int i=0; i < 3000; i++) {
            //
            fraction1.invert();
            // possibly some code here with random execution time.
            fraction1.report(id);
        };
        // sched_yield(); // absolute priority higher than 0
        // give up processor and allow another thread to use
        the processor
};

int main(int argc, char *argv[]) {
    /* Define behaviour for divide-by-zero */
    // SIGFPE is for integer divide-by-zero too.
    // Single Unix Specification defines SIGFPE as "Erroneous arithmetic operation."
    // https://stackoverflow.com/questions/6121623/catching-exception-divide-by-zero
    std::shared_ptr<void(int)> handler();
    signal(SIGFPE, [](int signum) {throw std::logic_error("FPE"); });
    [](__sighandler_t f) { signal(SIGFPE, f); });

    fraction1.report('P');

    thread threadA(threadFunc, 'A');
    thread threadB(threadFunc, 'B');
    mutex MyMutex; // bad, mutex as data member, make every class
    has mutex.
    // sleep(1);
    threadB.join();
    threadA.join();

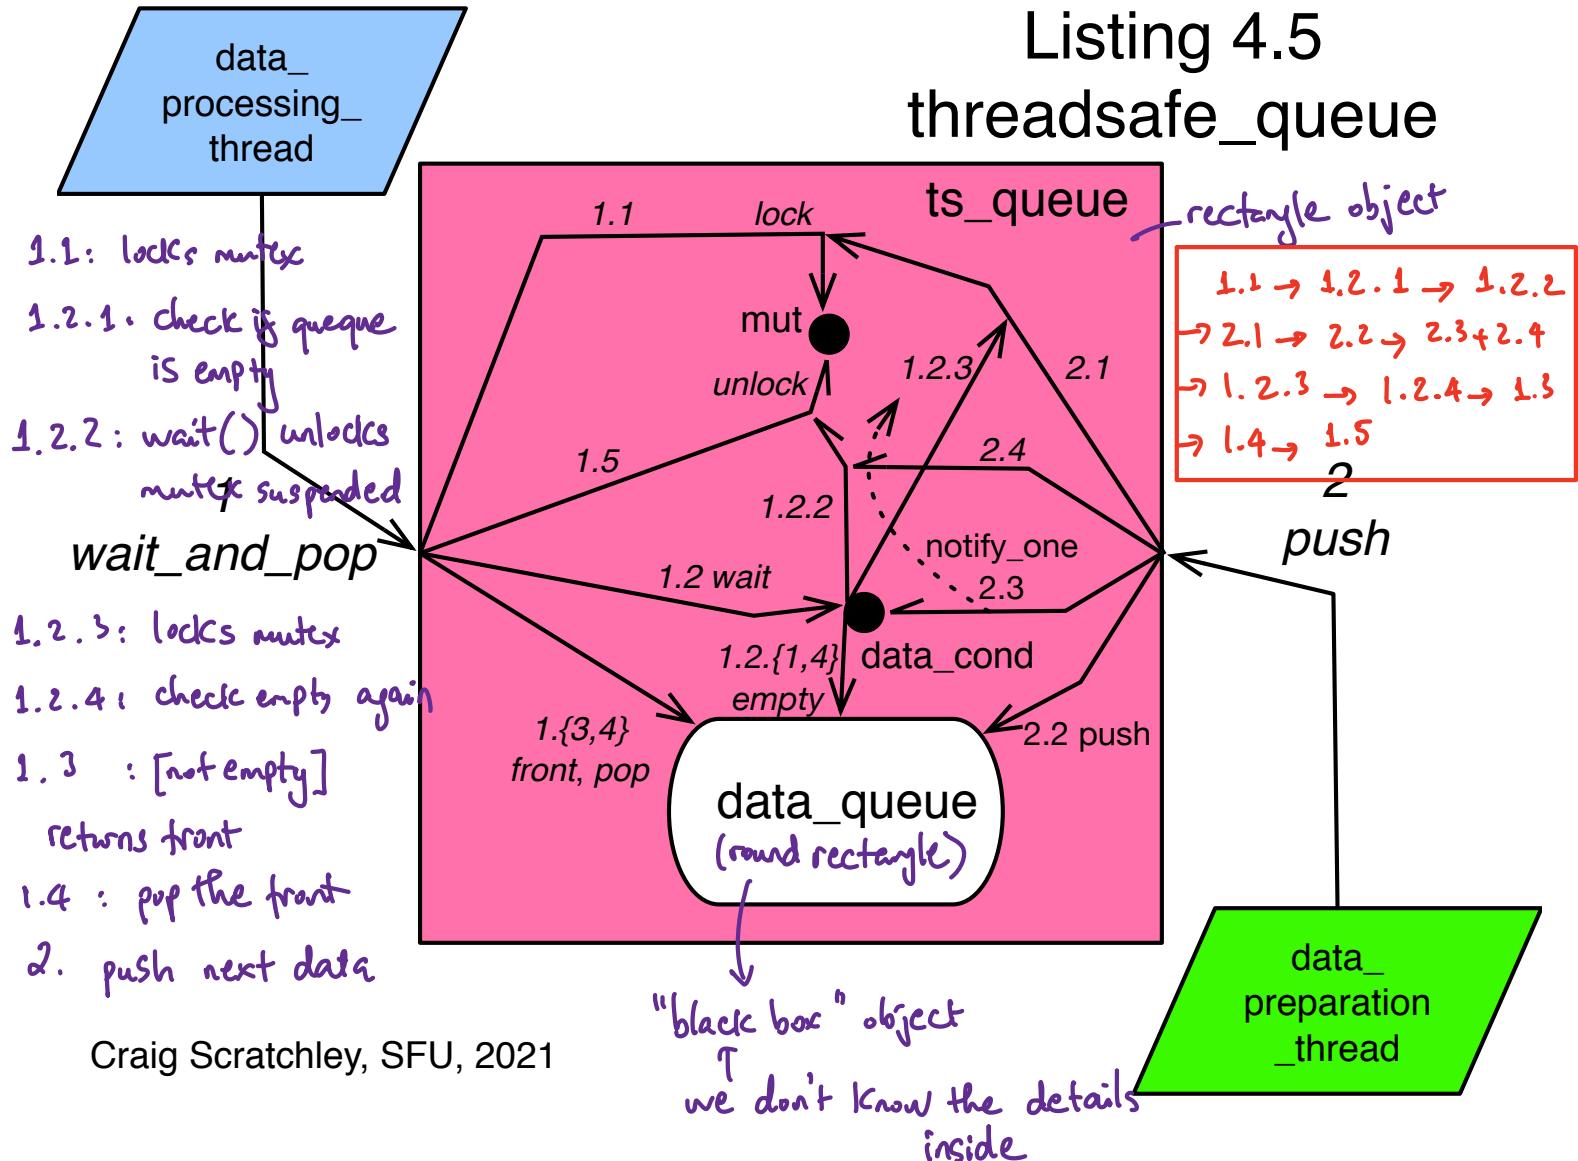
    cout << "Primary thread ending" << std::endl;
    return EXIT_SUCCESS;
}
//

```

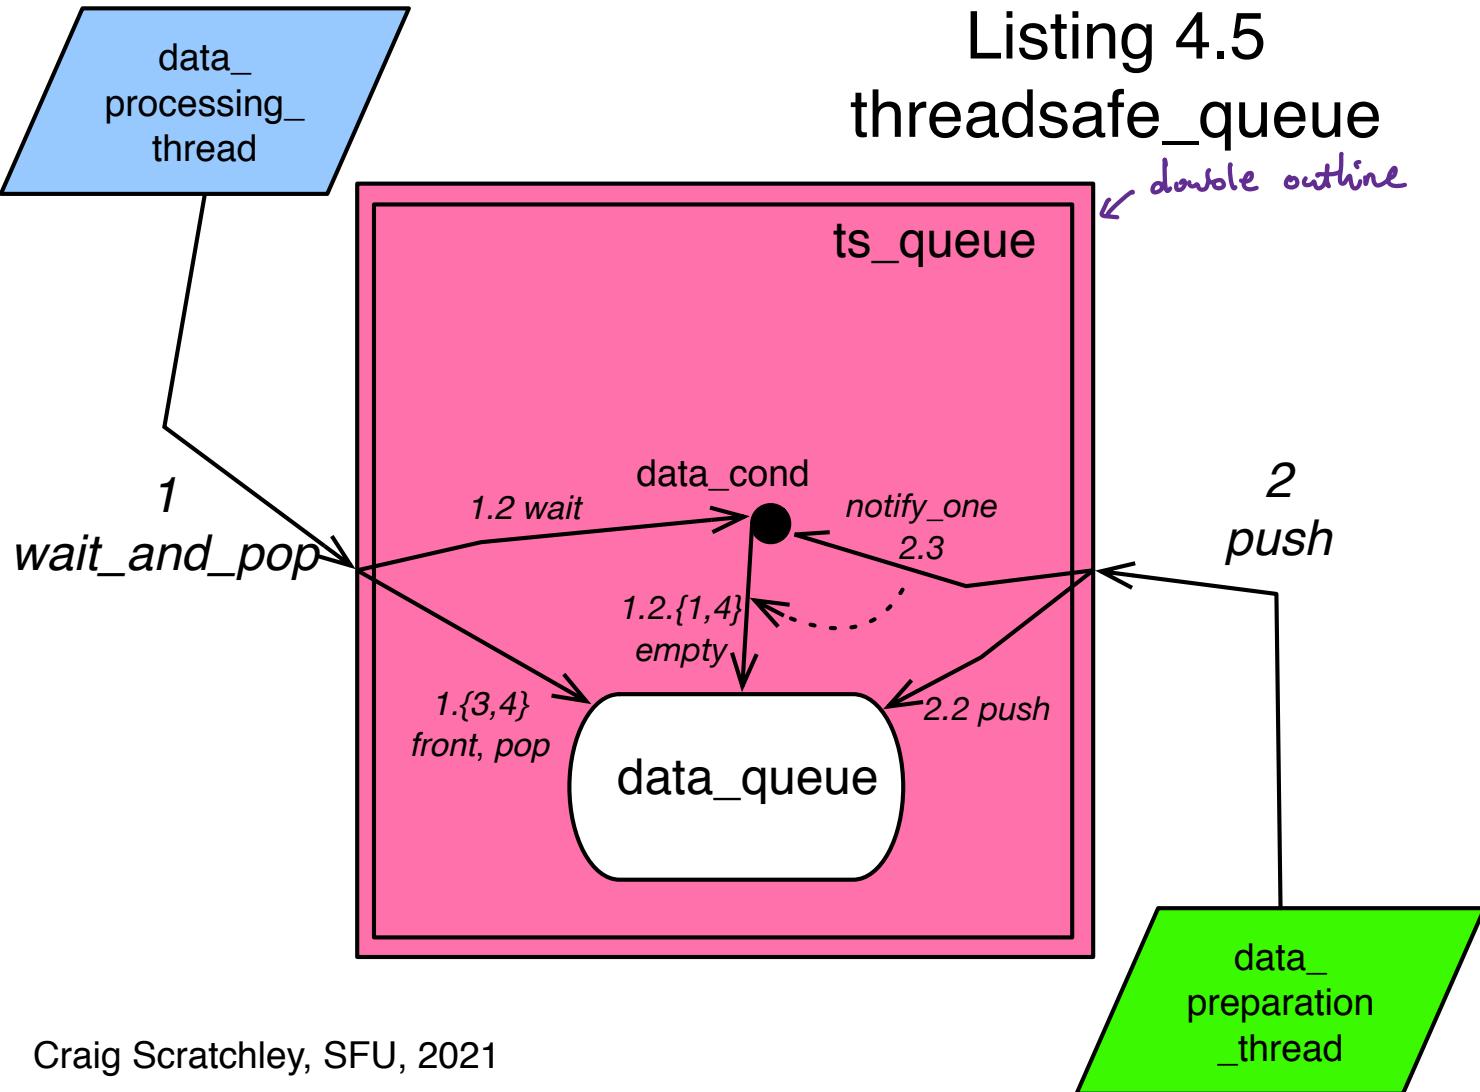
thread(file)

CXX → C++

Listing 4.5 threadsafe_queue



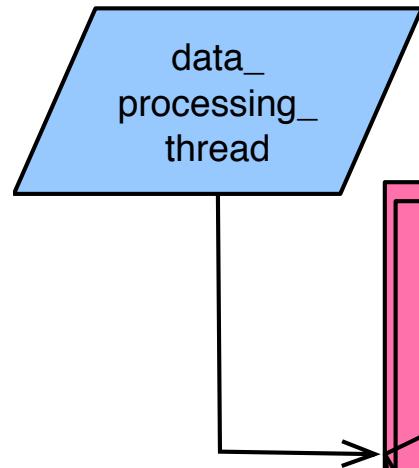
Listing 4.5 threadsafe_queue



threadsafe_queue w/ drain func

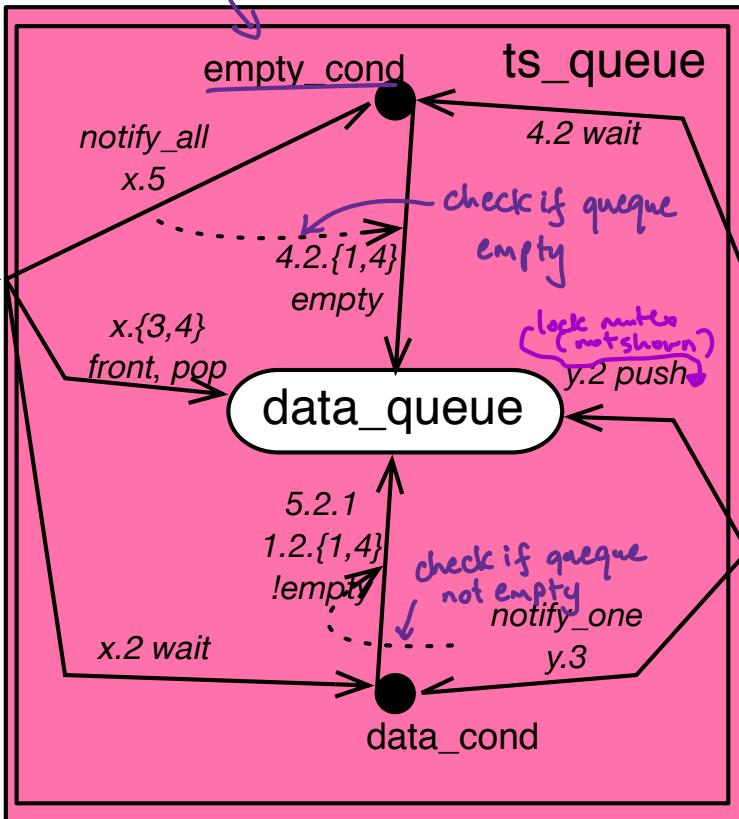
wait for data-queue becomes empty

threadsafe_queue with drain function



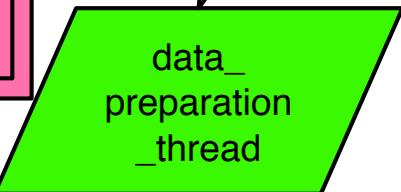
1,5
`wait_and_pop`

- 1.1 : locks mutex
- 1.2: call `wait()`
- 1.2.1 : check if !empty
- 1.3/4 : front, pop
- 2. 5 : `notify_all()`
- 2: `push()`



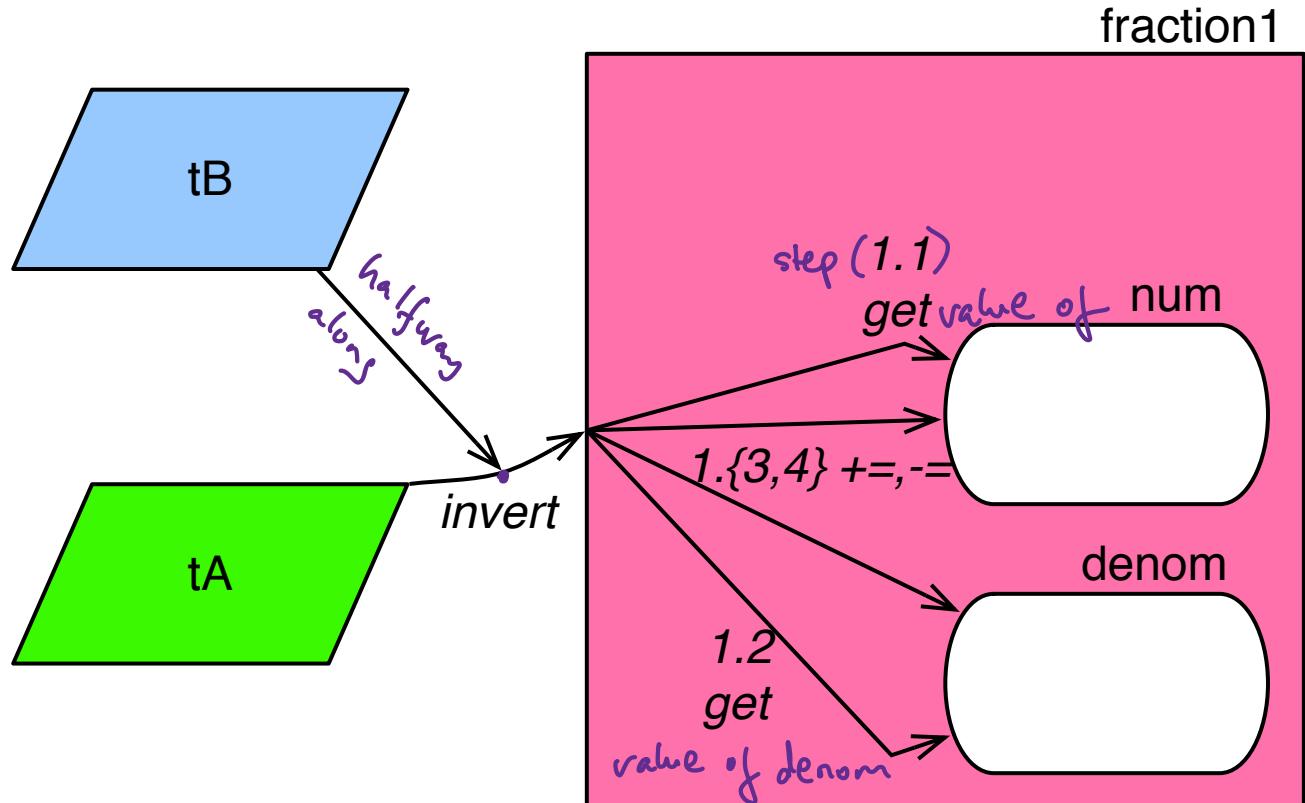
4. drain
4.2

4
drain
we know it's empty
we check and it is empty
→ we don't need to wait
if !empty, we drain data



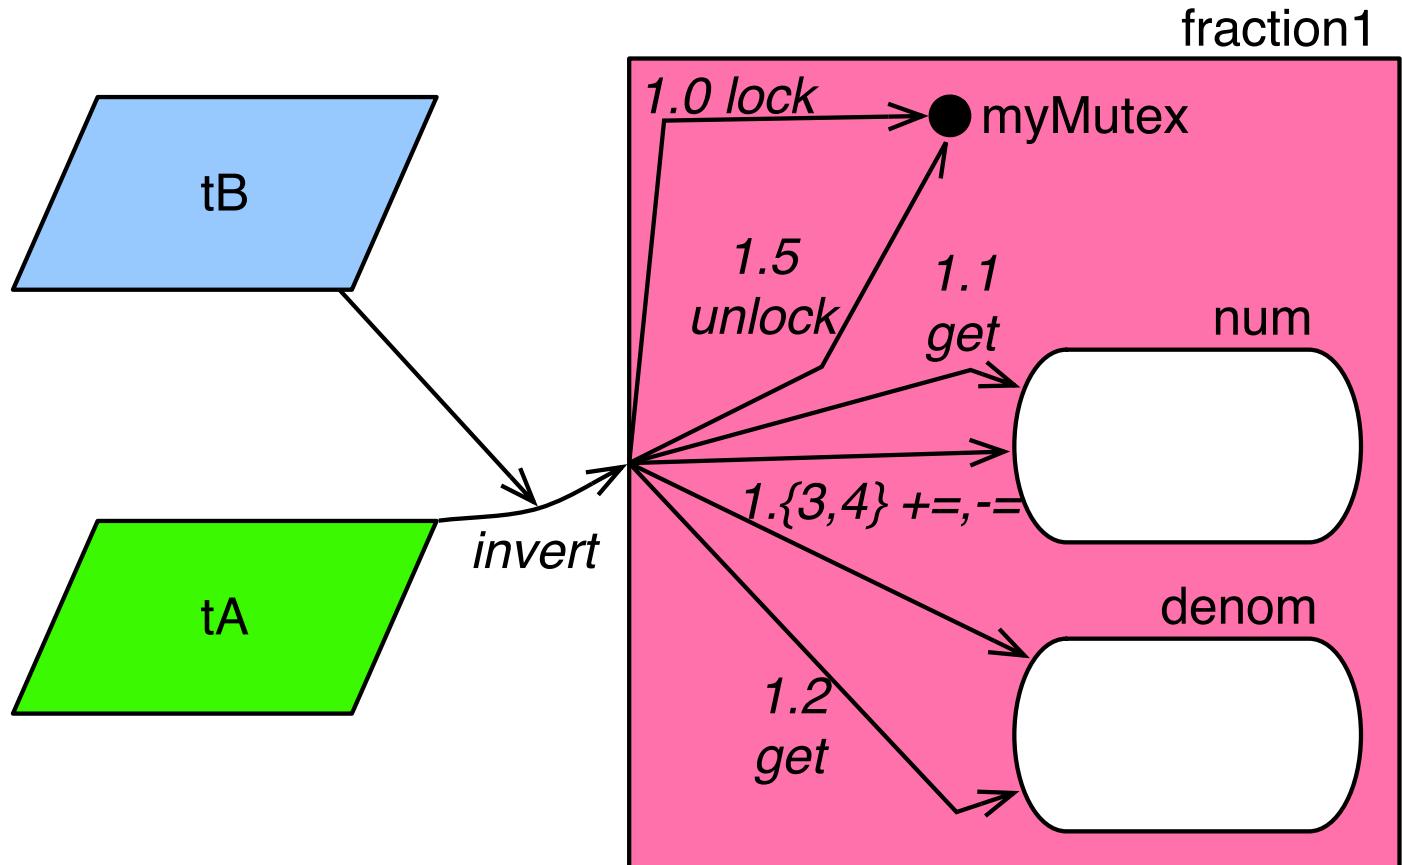
Craig Scratchley, SFU, 2020

two threads access a fraction object in an unsafe manner



Craig Scratchley, SFU, 2021

two threads access a fraction object protected by a mutex



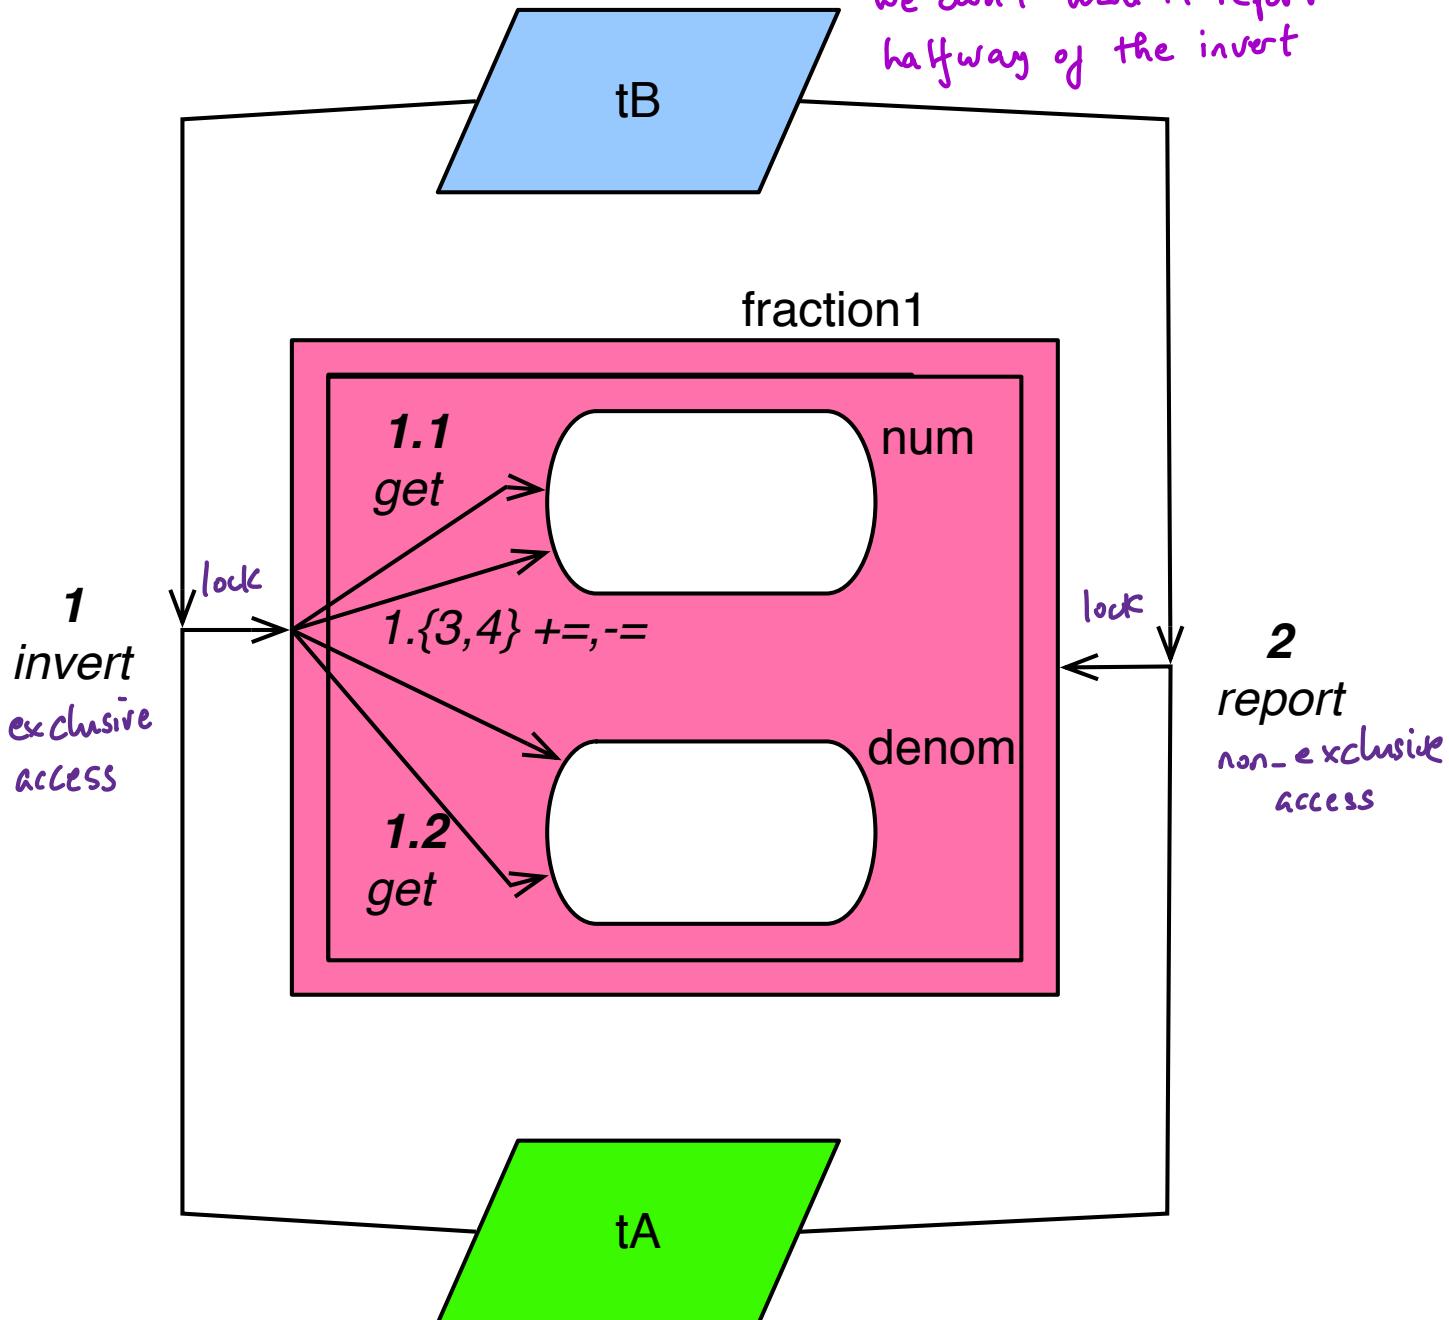
Craig Scratchley, SFU, 2020

two threads access a fraction object

in a protected manner

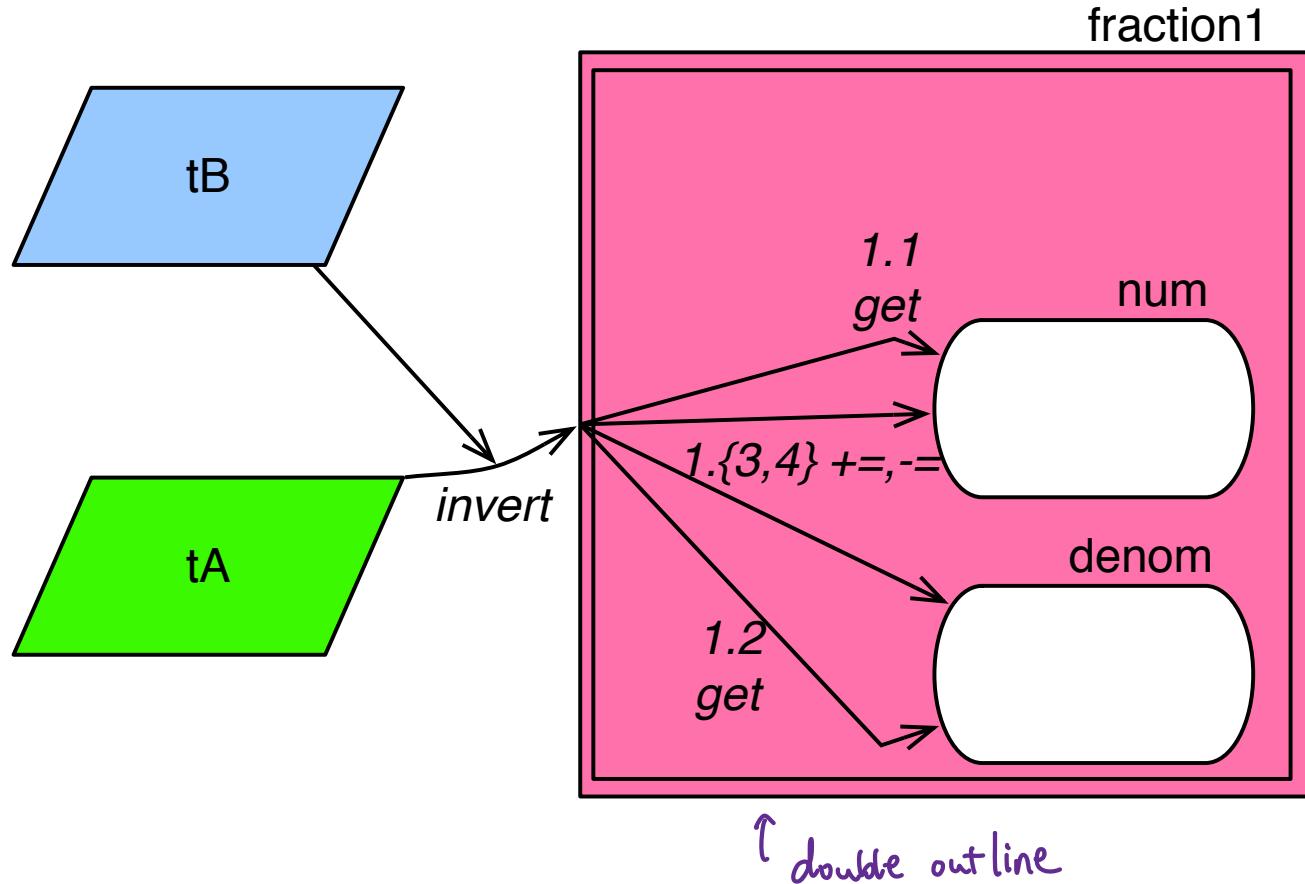
more report() called than invert() \rightarrow shared mutex

we don't want it report
halfway of the invert



Craig Scratchley, SFU, 2020

two threads access a fraction object in a protected manner

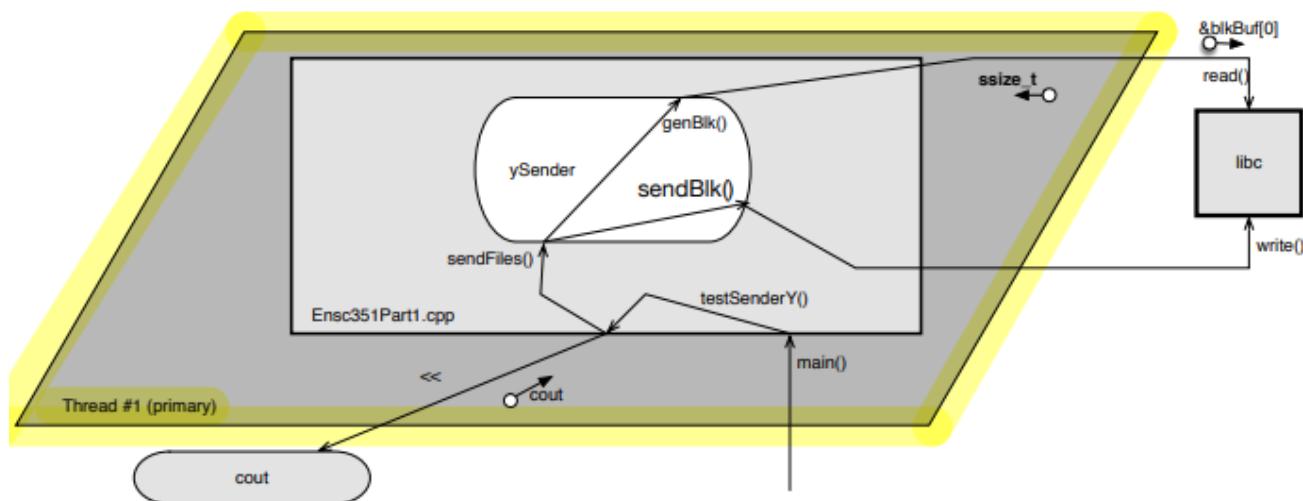


Arrow
YMODEM
CGN (Part1)

Figure 4. YMODEM Batch Transmission Session (2 files)

SENDER	RECEIVER	
"sending in batch mode etc."	"sb foo.c baz.c<CR>"	
SOH 00 FF foo.c NUL[123] CRC CRC	C (command:rb)	this class is associated with ——— this class
SOH 01 FE Data[128] CRC CRC	ACK	this class is dependent upon - - - → this class
SOH 02 FC Data[128] CRC CRC	C	this class inherits from → this class
SOH 03 FB Data[100] CPMEOF[28] CRC CRC	ACK	this class has ○ this interface
EOT	ACK	this class is a realisation of - - - ▷ this class
EOT	NAK	you can navigate from this class to → this class
SOH 00 FF baz.c NUL[123] CRC CRC	ACK	these classes compose without belonging to ◊ this class
SOH 01 FB Data[100] CPMEOF[28] CRC CRC	C	these classes compose and are contained by ◆ this class
EOT	ACK	this object sends a synchronous message to → this object
EOT	NAK	this object sends an asynchronous message to ➡ this object
SOH 00 FF NUL[128] CRC CRC	ACK	
	ACK	

A diagram for Part 1 of Multipart Project using Collaboration Graph Notation. The global object cout and primary thread have also been drawn.



C&N (part 2)
C&N (part 3)

Ensc351Part2.cpp

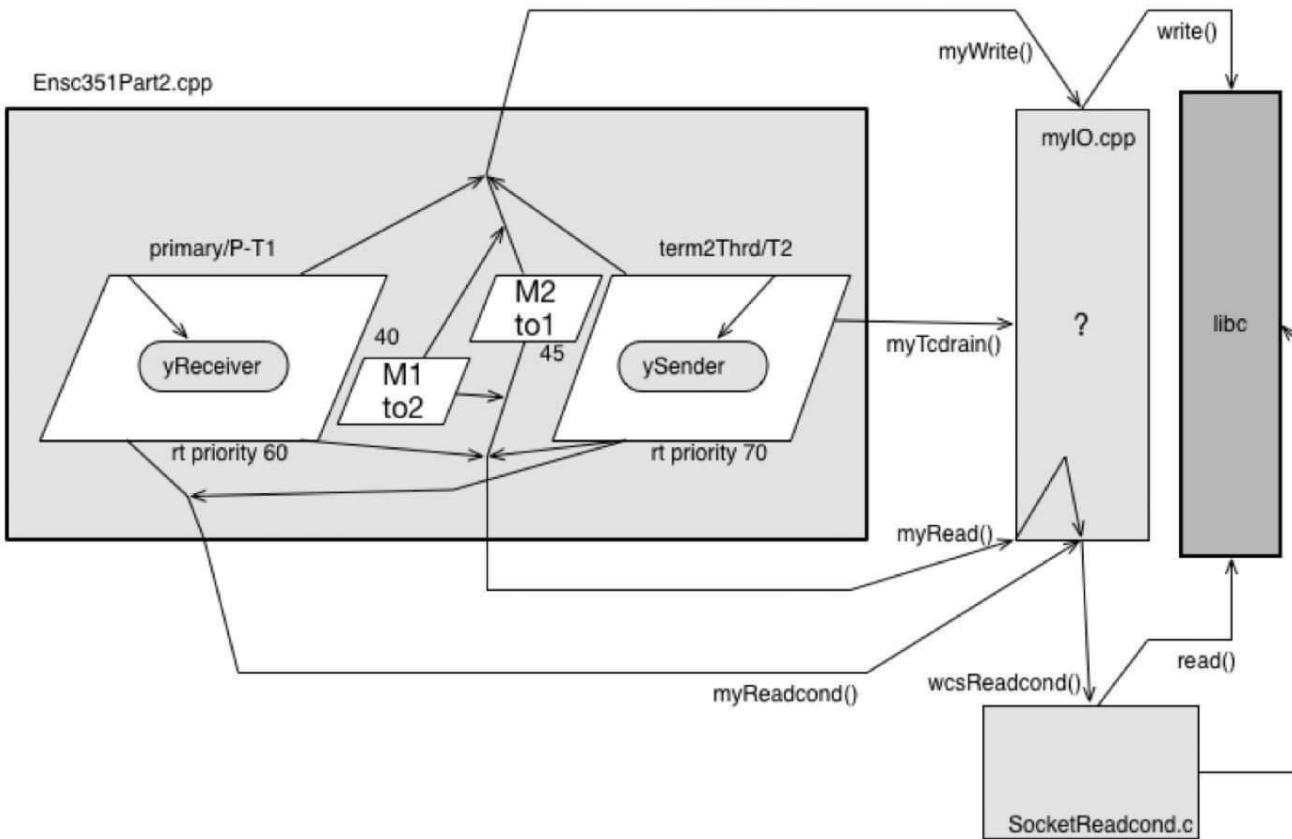
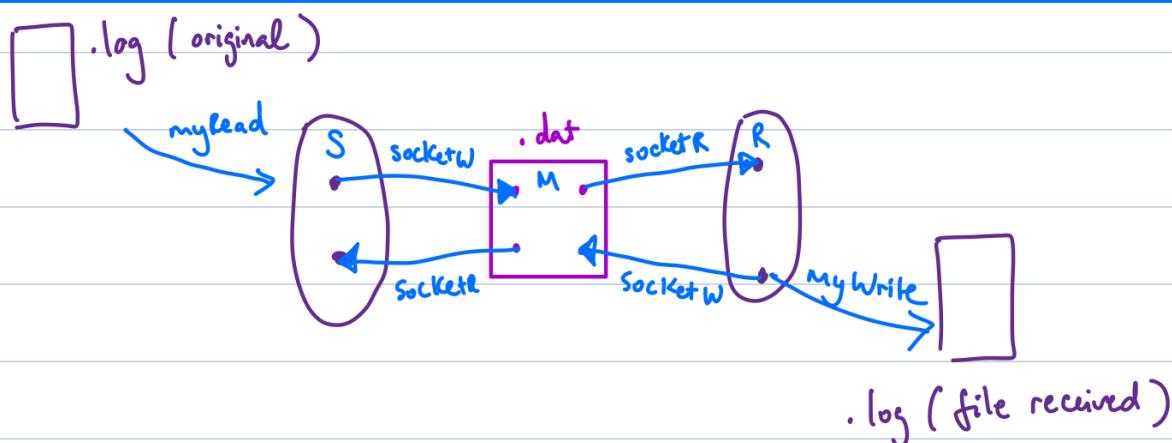
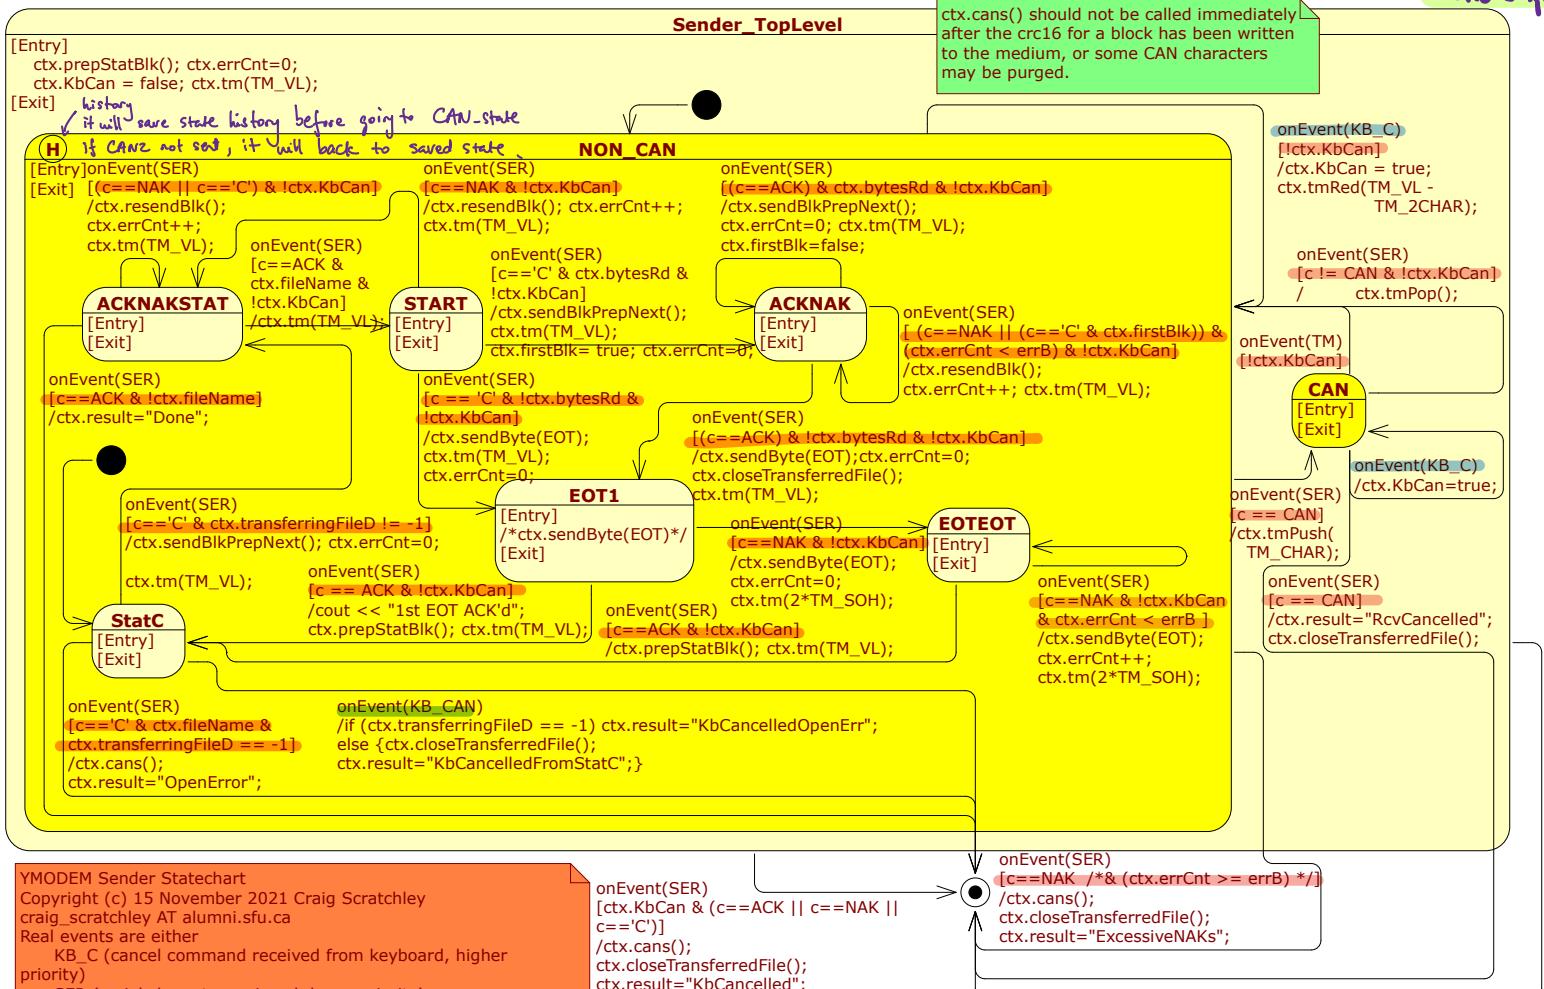


Figure 2: Collaboration Graph showing function calls for socketpair communication



pair	paired	A	B
-1		✓	
-2	-2	✓	✓



YMODEM Sender Statechart
 Copyright (c) 15 November 2021 Craig Scratchley
 craig.scratchley AT alumni.sfu.ca

Real events are either
 KB_C (cancel command received from keyboard, higher priority)
 SER (serial character received, lower priority),
 TM_VL (Very Long timeout) gives 60 seconds

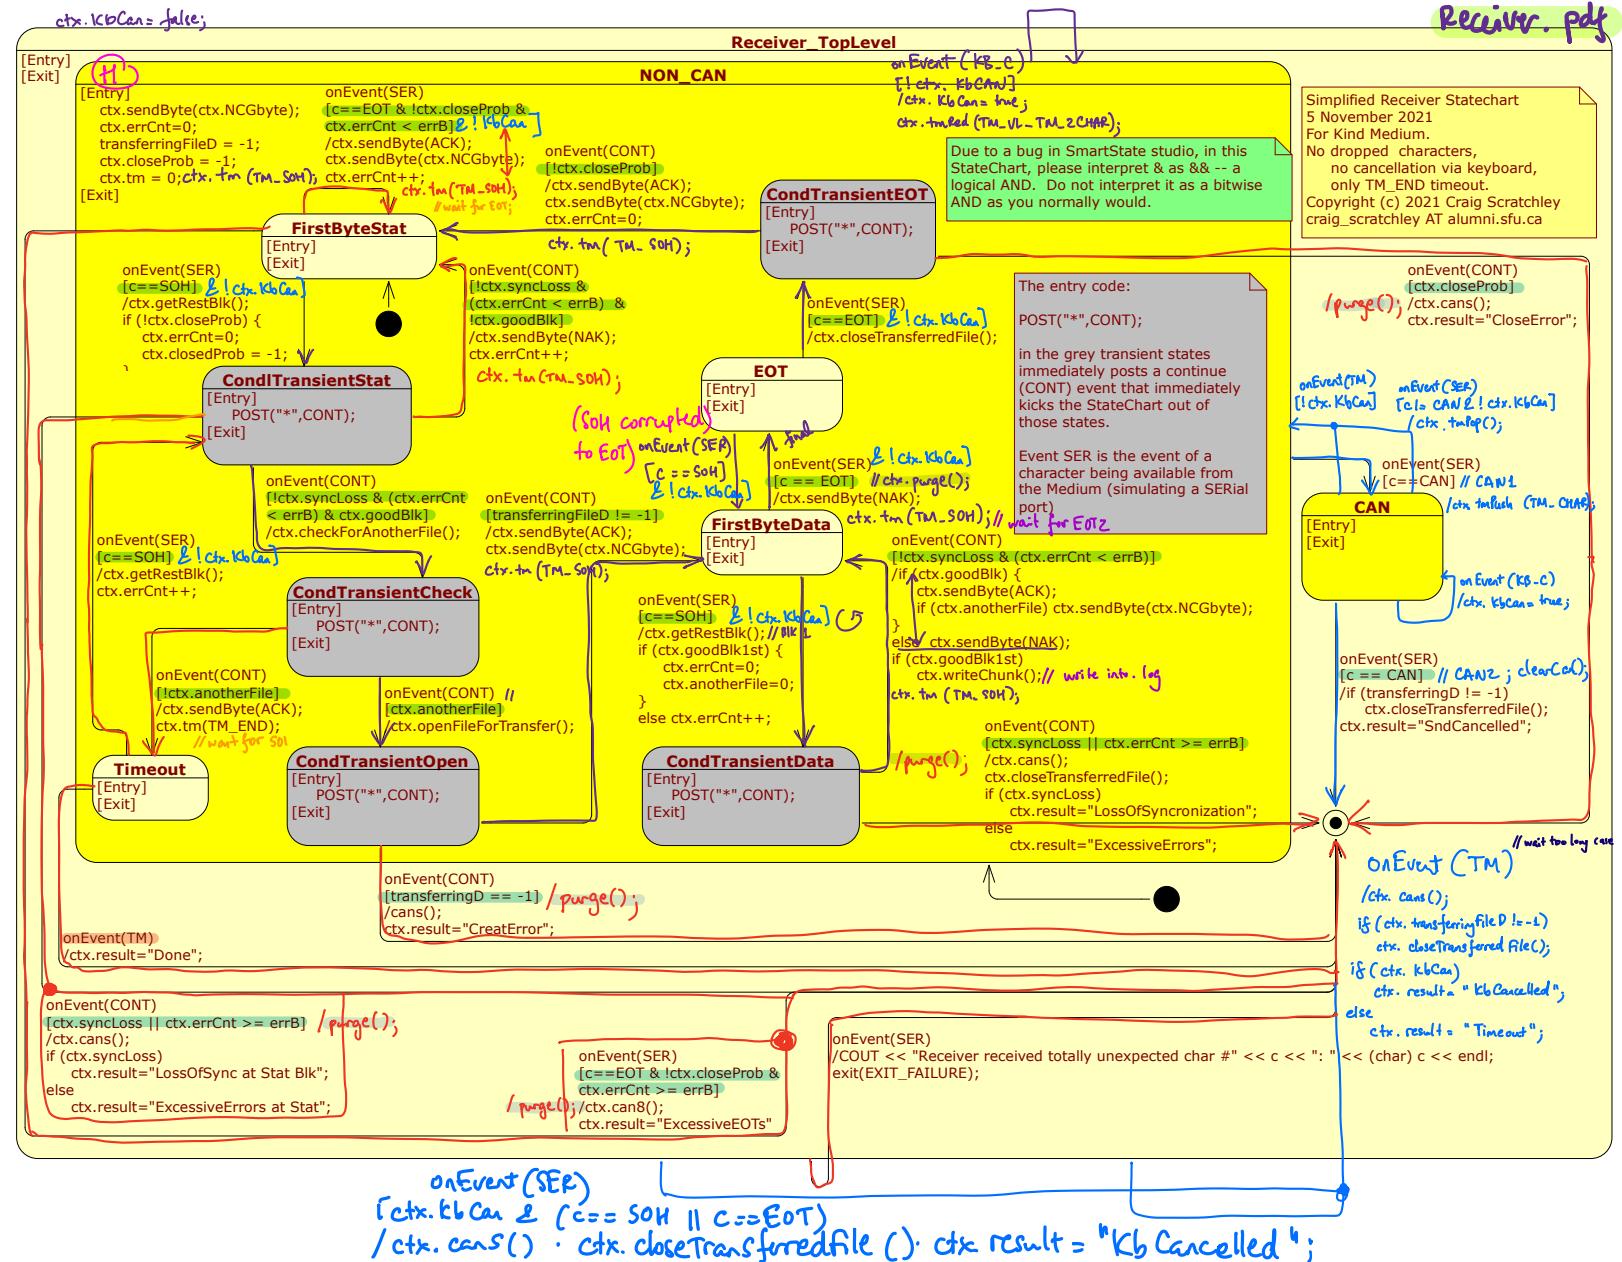
Unless a fast simulation has been chosen ...

TM_SOH (normal timeout waiting for SOH) gives 10 seconds
 TM_CHAR (inter-character timeout) gives 1 second
 TM_2CHAR gives a period longer than the
 inter-character timeout of 1 sec.

onEvent(SER)
 [ctx.KbCan & (c==ACK || c==NAK || c=='C')]
 /ctx.cans();
 ctx.closeTransferredFile();
 ctx.result="KbCancelled";

onEvent(SER)
 [c==NAK /*& (ctx.errCnt >= errB) */]
 /ctx.cans();
 ctx.closeTransferredFile();
 ctx.result="ExcessiveNAKs";

onEvent(TM)
 /ctx.cans();
 ctx.closeTransferredFile();
 if (ctx.KbCan)
 ctx.result="KbCancelled";
 else
 ctx.result="Timeout";



For reference, below is a description of the public member functions and variables in class *ReceiverY*, the receiver context (an object of which is referenced as *ctx* in the receiver state chart actions and guard conditions). If you need to add additional functions and/or variables to this class, the *SenderY* class, or the *PeerY* parent class, please check with Craig first.

Function or Variable	Description
void cans()	Send CAN_LEN copies of CAN characters in a row to the peer, to inform it of the cancelling of a session.
void getRestBlk()	<p>Call only after an SOH character has been received and posted to the receiver state chart. The function tries to read at least 132 more characters to form a complete block. The function will set or reset a Boolean variable, <i>goodBlk</i>. This variable will be made false if either</p> <ul style="list-style-type: none"> • 132 bytes have not yet been received and another byte does not arrive within 1 second of the last byte (or within 1 second of the function being called). • if a good copy of the block has not already been received, 132 bytes (or more) are received and the block completed using the first 132 received bytes has something wrong with it, like the <i>crc16</i> being incorrect. When it is released, see <i>getRestBlk()</i> in Part 2 solution for details. <p>The function will also set or reset another Boolean member variable, <i>syncLoss</i>. <i>syncLoss</i> will only be set to true when <i>goodBlk</i> is set to true AND there is a fatal loss of synchronization as described in the YMODEM specification. The first time each block is received and is good, <i>goodBlk1st</i> will be set to true. This is an indication of when the data in a block should be written to disk. If <i>goodBlk1st</i> is false, or in any case if more than 132 bytes were received, then the <i>purge()</i> function (see below) will be called before returning from <i>getRestBlk()</i>.</p>
void writeChunk()	Write the data in a received block to disk.
void clearCan()	Read and discard up to (CAN_LEN – 2) contiguous CAN characters. Read characters one-by-one in a loop until either nothing is received over a 2-second period or a character other than CAN is received. If received, send a non-CAN character to the console.
void purge()	The <i>purge()</i> subroutine will read and discard characters until nothing is received over a 1-second period.
int errCnt /*in PeerY*/	A variable that counts a sequence of 'C' or NAK characters sent, 10-second timeouts, or ACK characters sent due to repeated blocks. The initial 'C' does not add to the count. The reception of a good block for the first time resets the count (see <i>goodBlk1st</i> below).
int closeProb	0 indicates that there was no problem closing the file just transferred. A positive value is the errno encountered when trying to close the file. -1 indicates that the program has not yet tried to close the current file (or there is no current file).
bool goodBlk	A Boolean variable that indicates whether the last block received was good (or deemed good) or whether it had problems (and was not deemed to be good).

bool goodBlk1st	A Boolean variable that indicates that a good copy of a block being sent has been received for the first time. It is an indication that the data in the block can be written to disk.
bool syncLoss	A Boolean variable that indicates whether or not a fatal loss of synchronization has been detected.

The *ReceiverY* class, like the *SenderY* class described below, is a class derived from *PeerY*. Here is a description of the public member functions and variables you can use in *PeerY*:

Function or Variable	Description
void sendByte(uint8_t)	Send a byte to the remote peer across the medium
void tm(int tmSeconds)	set a timeout time at an absolute time <i>tmSeconds</i> into the future. That is, determine an absolute time to be used for the next YMODEM timeout by adding <i>tmSeconds</i> to the current time.
void tmPush(int tmSeconds)	Store the current absolute timeout, and create a temporary absolute timeout <i>tmSeconds</i> into the future.
void tmPop()	Discard the temporary absolute timeout and revert to the stored absolute timeout
void tmRed(int secsToReduce)	Make the absolute timeout time earlier by <i>secsToReduce</i> seconds.
std::string result	Points to a string giving the “result” when the session ends.
int transferringFileD	Descriptor for file being read from or written to.
int errCnt	See description in derived classes for interpretation for each class.

Note that the *errCnt* variable described in the derived classes actually lives here in *PeerY*.

As written above, if you need to add additional functions and/or variables to this base class, the *SenderY* class mentioned on the next page, or the *ReceiverY* class, please check with Craig first. A solution, as will be provided by Craig, uses only the provided functions and variables.

ReceiverY.
member func()
definition

For your information, here is a description of the public member functions and variables used in *SenderY*, the sender context class (an instance of which is referenced as *ctx* in the sender statechart actions and guard conditions):

Member Function or Variable	Description
void can8()	Send to the peer CAN_LEN copies of CAN characters in pairs, each pair separated by slightly over 1 second, to inform it of the canceling of a file transfer.
void sendBlkPrepNext()	Sends for the first time the block recently prepared (less the last byte), then updates the variable <i>bytesRd</i> and tries to prepare the next block. <i>bytesRd</i> is set to 0 and a block is not actually prepared if the input file is empty or the end of the input file was reached when the block recently prepared was filled in. Finally, after the rest of the block has been sent, dumps any received glitches and sends the last byte of block being sent (the last byte of crc16).
void resendBlk()	Resends the block that had been sent most recently or an empty zero-numbered “stat” block (less its last byte in both cases), and then dumps any received glitches and sends the last byte.
clearCan()	Read and discard contiguous CAN characters. Read up to (CAN_LEN – 2) characters one-by-one in a loop until either nothing is received over a 1-second period or a character other than CAN is received. If received, send a non-CAN character to the console.
int errCnt /*in PeerY*/	Counts the number of times that a block or EOT is resent.
ssize_t bytesRd	The number of bytes last read from the input file.
bool firstBlk	True if first block of file content data has not yet been acknowledged with an ACK.

Since we don't ask you, at this point, to test your Statechart by generating code from it and running the code, you can get full marks for your work as long as you have made a sincere attempt to handle in a reasonable way all the issues identified at the top of this document.

Since this part of the project will not require you to test your work, it should take less time than recent parts of the project.

```

#endif
using namespace std;
//Unnamed namespace
namespace{
    class socketInfoClass;
    typedef shared_ptr<socketInfoClass> socketInfoClassSp;
    map<int, socketInfoClassSp> desInfoMap {
        {-2, make_shared<socketInfoClass>(-1)}, // marker for a
descriptor whose pair is closed
        {0, nullptr}, // init for stdin,
        {1, nullptr}, //             stdout,
        {2, nullptr} //             stderr
    };

    // A shared mutex used to protect desInfoMap so only a
single thread can modify the map at a time.
    // This also means that only one call to functions like
mySocketpair() or myClose() can make progress at a time.
    // This mutex is also used to prevent a paired socket from
being closed at the beginning of a myWrite or myTcdrain function.
    shared_mutex mapMutex;

    class socketInfoClass {
        unsigned totalWritten = 0; // byte written to a buffer
        unsigned maxTotalCanRead = 0; // max total bytes can be read
        • condition_variable cvDrain; // by a thread in a socket
        • condition_variable cvRead;
    #ifndef WCSREADCOND
        • condition_variable cvRead;
    #endif
        #ifdef CIRCBUF
            CircBuf<char> circBuffer;
            bool connectionReset = false;
        #endif
        mutex socketInfoMutex;
    public:
        int pair; // Cannot be private because myWrite and
myTcdrain using it.
        // -1 when descriptor closed, -2 when paired
descriptor is closed
    };
}

// funcs below to unnamed namespace can only be used
in this file but not others. Keep it private!!
}

// avoid risk when combine code w/ partner whose also
have same name class. So only our class can access
this unnamed namespace.
myIO.cpp

```

*this map is not thread safe
→ need to lock mutex*

used to blk reading activity

(1)

(2)

(3)

```

int draining()
{ // operating on object for paired descriptor of original des
    unique_lock<mutex> socketLk(socketInfoMutex);

    // paired descriptor could have been closed before we
    constructed socketLk
    // once the reader decides the drainer should wakeup, it
    should wakeup
    if (pair >= 0 && totalWritten > maxTotalCanRead)
        cvDrain.wait(socketLk); // what about spurious wakeup?
        cvDrain.wait(socketLk, [this]{return pair < 0 || totalWritten <= maxTotalCanRead;});
    // these 2 condition code can replace w/ comment line
    if (pair == -2) {
        errno = EBADF; // check errno
        return -1;
    }
    return 0;
}

```

i.e.: If one lock in thread 2, one in thread 3 if you can't lock one, the other also can't you can't lock each other, but need another thread 4 to lock them

```

unique_lock<mutex> socketLk(socketInfoMutex,
defer_lock);
unique_lock<mutex> condPairLk(des_pair->socketInfoMutex, defer_lock);
lock(condPairLk, socketLk); // lock both mutex at the same time,
pair = -1; // this is first socket in the pair to be
closed
int mySocketpair(int domain, int type, int protocol, int des[2])
{
    lock_guard<shared_mutex> desInfoLk(mapMutex);
    int returnVal = socketpair(domain, type, protocol, des);
    if(returnVal != -1) {
        move inside if() what happen if we switch the order of these 2 lines?
        socketpair is thread safe (which not modify
        our Map, so it can be put above mutex → holds
        mutex for shorter time → more efficient so
    }
}
```

```

        desInfoMap.emplace(des[0], other thread don't wait too long
make_shared<socketInfoClass>(des[1]));
        desInfoMap.emplace(des[1],
make_shared<socketInfoClass>(des[0]));
    }
    return returnVal;
}

```

// check if put mutex at
the beginning gives any error?

tw: total written
mTCR: maximum total can be read

threadable part 3 sol

Order of execution of code in Ensc351Part3-test.cpp and some output and details of the solution

threadT41 (priority 50)	threadT32 (priority 70 to 40 to 80 to 40)	threadT42 (priority 60)	daSktPr[0] - 3		daSktPr[1] - 4	
			tW/mTCR	pair	cvDrn	tW/mTCR
			-	-	-	-
	mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr); mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr1); mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr2); PE_NOT(myWrite(daSktPr[0], "abcd", 4), 4); // write 4 chars to socket 0 posixThread threadT42([=, threadT42func]; /* X */ myTcdrain(daSktPr[0])); // blocked	socket 0	0/0	4	0/0	3 /
read 4 char from thread T32 0 ↓ return to Tedrain			0/0	4	4/0	3 /
/* A */ /myReadcond(daSktPr[1], Ba, 20, 12, 0, 0); // blocked		socket 1	0/0	4	4/0	3 T32/
// take 8 more bytes (read 4 bytes) unblock	setSchedPrio(40); /* finished: statement A: result was 14 Ba: abcd123456789 */ myReadcond(daSktPr[1], Ba, 20, 0, 0); // returned 0; myWrite(daSktPr[1], "Will not be read", 17); /* B */ /myReadcond(daSktPr[1], Ba, 20, 12, 0, 0); // blocked	/* finished: statement A: result was 14 Ba: abcd123456789 */ PE_NOT(myWrite(daSktPr[0], "123456789", 10), 10); // write 10 more bytes // close socket 0, unblock 2 threads (include NULL) in socket 2	5/0	4	4/0	3 T32/
↑ ↑ max min unblock don't bore me ↑ till you get 12bytes	setSchedPrio(80); PE_NOT(myWrite(daSktPr[0], "xyz", 4), 4); PE(myTcdrain(daSktPr[0])); PE(myClose(daSktPr[0])); // returned 0 setSchedPrio(40);	/* finished: statement B: result was 4 Ba: xyz */ /* we closed, so only 4 byte read.	5/0	4	T42	4/0
/* C */ /myReadcond(daSktPr[1], Ba, 20, 1, 0, 0); // bl'ked	PE(myClose(daSktPr[0]));	/* finished: statement X: result was 0 */ myTcdrain(daSktPr[1]); threadT41.join();	22/0	4	14/20	3 /
/* finished: C: res was -1 errno 104: Connection reset by peer */ myReadcond(daSktPr[1], Ba, 20, 1, 0, 0); // will return 0 myWrite(daSktPr[1], "Added", 6); // ret -1 errno 32: Broken pipe /* D */ /myRead(daSktPr[2][1], Ba, 20); // blocked	PE(myClose(daSktPr1[0]));	/* error due to my write "will not be read." above	22/0	4	T42	4/20
/* finished: statement D: result was 4 Ba: mno */ /* E */ /myRead(daSktPr[2][1], Ba, 20);	PE(myClose(daSktPr2[0]));	/* finished: statement E: result was 0 */ myClose(daSktPr2[1]); // returned 0 myClose(daSktPr1[1]); // returned 0 myClose(daSktPr[1]); // returned 0 myClose(daSktPr[1]); // ret -1 errno 9: Bad file descriptor myRead(daSktPr[1], Ba, 20); // ret -1 errno 9 // ... /* end of threadT41 */	22/0	-1	4/20	-2 /
		/* end of threadT42 */				
		threadT42.join();				

```
/* Ensc351Part3-test.cpp -- October 21 -- Copyright 2021 Craig Scratchley */  
  
/* This program can be used to test your changes to myIO.cpp  
*  
* Put this project in the same workspace as your Ensc351 library project,  
* and build it.  
*  
* With 3 created threads for a total of 4 threads, the output that I get  
with my solution  
* is in the file "output-fromSolution".  
*  
*/  
  
#include <sys/socket.h>  
#include <stdlib.h> // for exit()  
#include <sched.h>  
#include "posixThread.hpp"  
#include "VNPE.h"  
#include "myIO.h"  
  
// #define OUTPUT_TO_FILE  
  
#ifdef OUTPUT_TO_FILE  
    #include <fstream>  
    std::ofstream COUT("output");  
    // When RUNning the Release configuration, output will appear in the  
Eclipse project.  
    // When using the "viaConnection" launcher, output will appear at path  
/home/osboxes/output  
#else  
    #include "AtomicCOUT.h"  
#endif  
  
  
#define REPORT0(S) COUT << threadName << ":" << #S << "; statement will now  
be started\n"; \  
    S; \  
    COUT << threadName << ":" << #S << "; statement has now finished\n";  
#define REPORT1(FC) {COUT << threadName << ":" << #FC << " will now be  
called\n"; \  
    int RV = FC; \  
    COUT << threadName << ":" << #FC << " result was " << RV; \  
    if (RV == -1) COUT << " errno " << errno << ":" << strerror(errno); \  
    COUT << "\n"; }  
#define REPORT2(FC) {COUT << threadName << ":" << #FC << " will now be  
called\n"; \  
    int RV = FC; \  
    COUT << threadName << ":" << #FC << " result was " << RV; \  
    if (RV == -1) COUT << " errno " << errno << ":" << strerror(errno); \  
    else if (RV > 0) COUT << " Ba: " << Ba; \  
    COUT << "\n"; }  
  
using namespace std;  
using namespace pthreadSupport;  
  
static int daSktPr[2]; // Descriptor Array for Socket Pair  
static int daSktPr2[2]; // Descriptor Array for 2nd Socket Pair
```

```
static int daSktPr1[2];      // Descriptor Array for another Socket Pair
cpu_set_t cpu_set;
int myCpu=0;

void threadT41Func(void) // starts at priority 50
{
    char     Ba[20];
    const char* threadName = "T41";
    PE_0(pthread_setname_np(pthread_self(), threadName));
    //
    // Blank lines below indicate that the statement above the blank line
    // will finish after one or more other threads in the process have made
    progress.
    REPORT2(myReadcond(daSktPr[1], Ba, 20, 12, 0, 0)); // will block until
myWrite of 10 characters

    REPORT2(myReadcond(daSktPr[1], Ba, 20, 0, 0, 0));
    REPORT1(myWrite(daSktPr[1], "Will not be read", 17));
    REPORT2(myReadcond(daSktPr[1], Ba, 20, 12, 0, 0)); // will block until
myClose(daSktPr[0])

    REPORT2(myReadcond(daSktPr[1], Ba, 20, 1, 0, 0)); // will return -1
with error 104, Connection reset by peer
    REPORT1(myWrite(daSktPr[1], "Will not be read", 17));
    REPORT2(myReadcond(daSktPr[1], Ba, 20, 1, 0, 0)); // will block until
myClose(daSktPr[0])

    REPORT2(myReadcond(daSktPr1[1], Ba, 20, 1, 0, 0)); // will return 0
    REPORT1(myWrite(daSktPr[1], "Added", 6));
    REPORT2(myRead(daSktPr2[1], Ba, 20));

    REPORT2(myRead(daSktPr2[1], Ba, 20));

    REPORT1(myClose(daSktPr2[1]));
    REPORT1(myClose(daSktPr1[1]));
    REPORT1(myClose(daSktPr[1]));
    REPORT1(myClose(daSktPr[1]));
    REPORT2(myRead(daSktPr[1], Ba, 20));
    REPORT2(myReadcond(daSktPr[1], Ba, 20, 0, 0, 0));
    REPORT1(myWrite(daSktPr[1], Ba, 20));
    REPORT1(myTcdrain(daSktPr[1]));
}

void threadT42Func(void) // starts at priority 60
{
    const char* threadName = "T42";
    PE_0(pthread_setname_np(pthread_self(), threadName));
    //
    REPORT1(PE_NOT(myWrite(daSktPr[1], "ijkl", 5), 5));
    REPORT0(posixThread threadT41(50, threadT41Func));
    REPORT1(myTcdrain(daSktPr[1])); // will block until myClose(daSktPr[0])

    REPORT1(myTcdrain(daSktPr[1]));
    REPORT0(threadT41.join());

    // output happens at this time from the above REPORT0
}
```

```
void threadT32Func(void) // priority 70 -> priority 40 -> priority 80 ->
priority 40
{
    const char* threadName = "T32";
    PE_0(pthread_setname_np(pthread_self(), threadName));
    //
    REPORT1(mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr));
    REPORT1(mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr1));
    REPORT1(mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr2));
    //
    REPORT1(PE_NOT(myWrite(daSktPr[0], "abcd", 4), 4));
    REPORT0(posixThread threadT42(60, threadT42Func));
    REPORT1(myTcdrain(daSktPr[0])); // will block until 1st myReadcond(...,
12, ...);

    REPORT1(setSchedPrio(40));
    REPORT1(PE_NOT(myWrite(daSktPr[0], "123456789", 10), 10)); // don't
forget nul termination character

    REPORT1(setSchedPrio(80));
    REPORT1(PE_NOT(myWrite(daSktPr[0], "xyz", 4), 4));
    REPORT1(PE(myTcdrain(daSktPr[0])));
    REPORT1(myClose(daSktPr[0]));
    REPORT1(setSchedPrio(40));

    sched_yield(); // I noticed once the above priority change was not
immediately respected
    REPORT1(myClose(daSktPr1[0]));

    REPORT1(PE_NOT(myWrite(daSktPr2[0], "mno", 4), 4));
    REPORT1(myClose(daSktPr2[0]));

    REPORT0(threadT42.join());
}

int main() {
    CPU_SET(myCpu, &cpu_set);
    const char* threadName = "Pri";
    PE_0(pthread_setname_np(pthread_self(), threadName));
    PE(sched_setaffinity(0, sizeof(cpu_set), &cpu_set)); // set processor
affinity for current thread

    // Pre-allocate some memory for the process.
    // Seems to make priorities work better, at least
    // when using gdb.
    // For some programs, the amount of memory allocated
    // here should perhaps be greater.
    void* ptr = malloc(20000);
    free(ptr);

    try{
        sched_param sch;
        int policy = -1;
        int primaryPriority = 90;
```

```

        getSchedParam(&policy, &sch);
        if (sch._sched_priority < 98)
            cout << "**** If you are debugging, debugger is not running at a
high priority. ****\n" <<
                " **** This could cause problems with debugging.
Consider debugging\n" <<
                    " **** with the proper debug launch configuration
****" << std::endl;
        cout << "Primary Thread was executing at policy " << policy << " and
priority " << sch.sched_priority << endl;
        sch._sched_priority = primaryPriority;
        setSchedParam(SCHED_FIFO, sch); //SCHED_FIFO == 1, SCHED_RR == 2
        getSchedParam(&policy, &sch);
        cout << "Primary Thread now executing at policy (should be 1) " <<
policy << " and priority (should be " << primaryPriority << ") " <<
sch.sched_priority << endl;

        REPORT0(posixThread T32(SCHED_FIFO, 70, threadT32Func));
        REPORT0(T32.join());

        return 0;
    }
    catch (std::system_error& error){
        cout << "Error: " << error.code() << " - " << error.what() << endl;
        return error.code().value();
    }
    catch (...) { throw; }
}

```

```

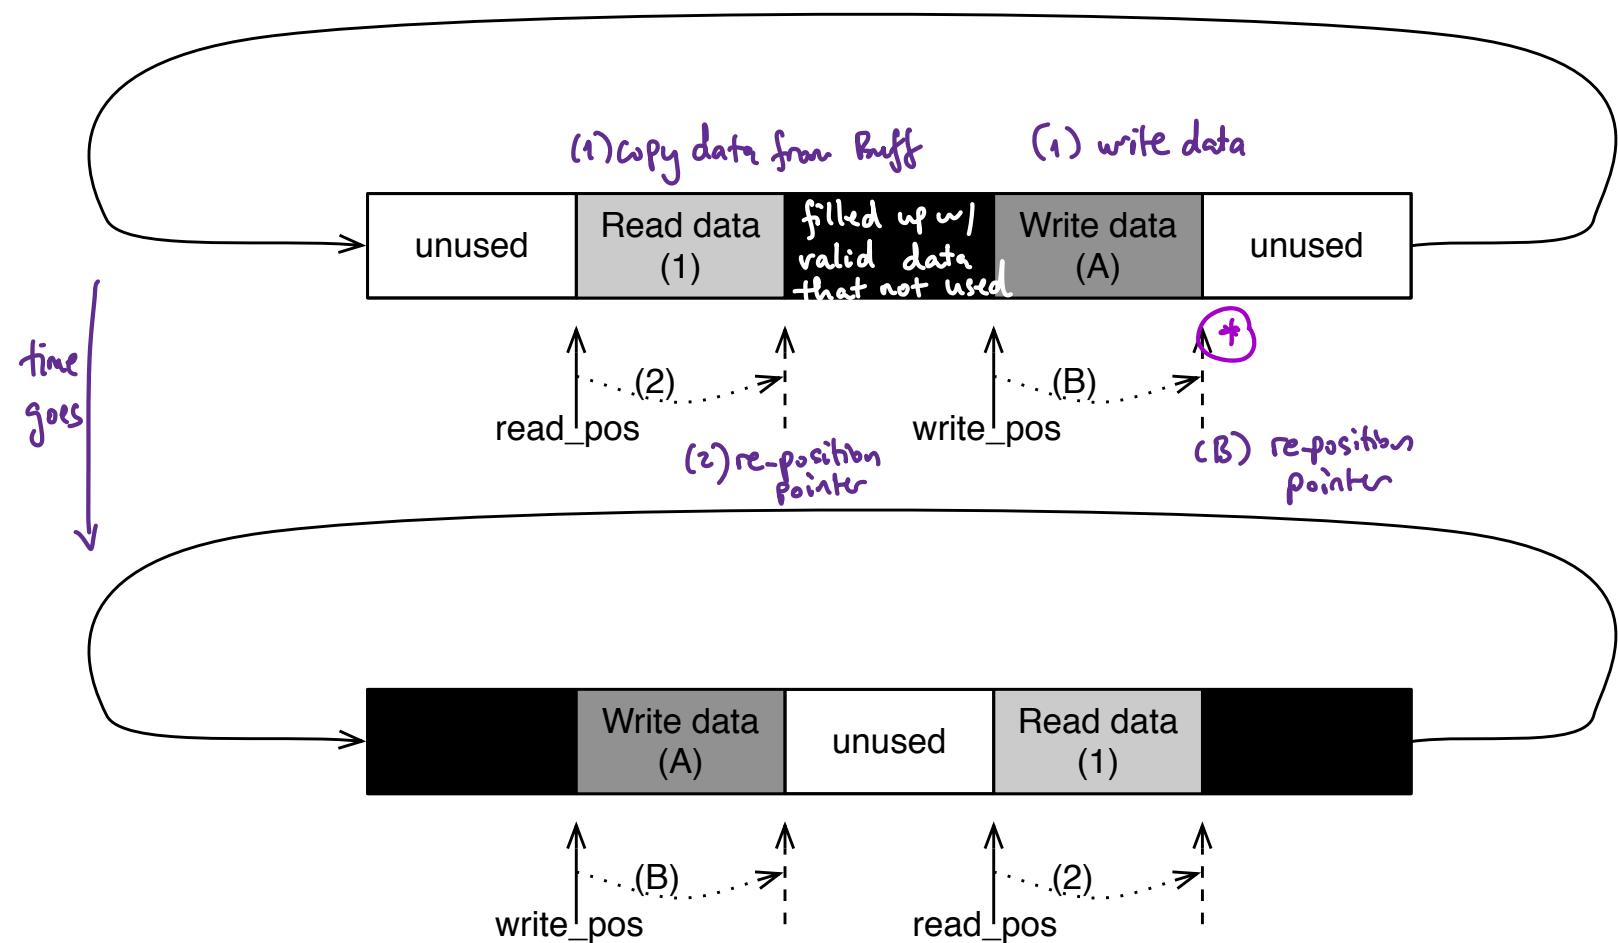
Pri: posixThread T32(SCHED_FIFO, 70, threadT32Func); statement will now be started
Pri: posixThread T32(SCHED_FIFO, 70, threadT32Func); statement has now finished
Pri: T32.join(); statement will now be started
T32: mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr) will now be called
T32: mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr) result was 0
T32: mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr1) will now be called
T32: mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr1) result was 0
T32: mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr2) will now be called
T32: mySocketpair(AF_LOCAL, SOCK_STREAM, 0, daSktPr2) result was 0
T32: PE_NOT(myWrite(daSktPr[0], "abcd", 4), 4) will now be called
T32: PE_NOT(myWrite(daSktPr[0], "abcd", 4), 4) result was 4
T32: posixThread threadT42(60, threadT42Func); statement will now be started
T32: posixThread threadT42(60, threadT42Func); statement has now finished
T32: myTcdrain(daSktPr[0]) will now be called
T42: PE_NOT(myWrite(daSktPr[1], "ijkl", 5), 5) will now be called
T42: PE_NOT(myWrite(daSktPr[1], "ijkl", 5), 5) result was 5
T42: posixThread threadT41(50, threadT41Func); statement will now be started
T42: posixThread threadT41(50, threadT41Func); statement has now finished
T42: myTcdrain(daSktPr[1]) will now be called
T41: myReadcond(daSktPr[1], Ba, 20, 12, 0, 0) will now be called
T32: myTcdrain(daSktPr[0]) result was 0
T32: setSchedPrio(40) will now be called
T32: setSchedPrio(40) result was 0
T32: PE_NOT(myWrite(daSktPr[0], "123456789", 10), 10) will now be called
T41: myReadcond(daSktPr[1], Ba, 20, 12, 0, 0) result was 14 Ba: abcd123456789
T41: myReadcond(daSktPr[1], Ba, 20, 0, 0, 0) will now be called
T41: myReadcond(daSktPr[1], Ba, 20, 0, 0, 0) result was 0
T41: myWrite(daSktPr[1], "Will not be read", 17) will now be called

```

```
T41: myWrite(daSktPr[1], "Will not be read", 17) result was 17
T41: myReadcond(daSktPr[1], Ba, 20, 12, 0, 0) will now be called
T32: PE_NOT(myWrite(daSktPr[0], "123456789", 10), 10) result was 10
T32: setSchedPrio(80) will now be called
T32: setSchedPrio(80) result was 0
T32: PE_NOT(myWrite(daSktPr[0], "xyz", 4), 4) will now be called
T32: PE_NOT(myWrite(daSktPr[0], "xyz", 4), 4) result was 4
T32: PE(myTcdrain(daSktPr[0])) will now be called
T32: PE(myTcdrain(daSktPr[0])) result was 0
T32: myClose(daSktPr[0]) will now be called
T32: myClose(daSktPr[0]) result was 0
T32: setSchedPrio(40) will now be called
T42: myTcdrain(daSktPr[1]) result was 0
T42: myTcdrain(daSktPr[1]) will now be called
T42: myTcdrain(daSktPr[1]) result was 0
T42: threadT41.join(); statement will now be started
T41: myReadcond(daSktPr[1], Ba, 20, 12, 0, 0) result was 4 Ba: xyz
T41: myReadcond(daSktPr[1], Ba, 20, 1, 0, 0) will now be called
T41: myReadcond(daSktPr[1], Ba, 20, 1, 0, 0) result was -1 errno 104: Connection
reset by peer
T41: myWrite(daSktPr1[1], "Will not be read", 17) will now be called
T41: myWrite(daSktPr1[1], "Will not be read", 17) result was 17
T41: myReadcond(daSktPr1[1], Ba, 20, 1, 0, 0) will now be called
T32: setSchedPrio(40) result was 0
T32: myClose(daSktPr1[0]) will now be called
T41: myReadcond(daSktPr1[1], Ba, 20, 1, 0, 0) result was -1 errno 104: Connection
reset by peer
T41: myReadcond(daSktPr1[1], Ba, 20, 1, 0, 0) will now be called
T41: myReadcond(daSktPr1[1], Ba, 20, 1, 0, 0) result was 0
T41: myWrite(daSktPr[1], "Added", 6) will now be called
T41: myWrite(daSktPr[1], "Added", 6) result was -1 errno 32: Broken pipe
T41: myRead(daSktPr2[1], Ba, 20) will now be called
T32: myClose(daSktPr1[0]) result was 0
T32: PE_NOT(myWrite(daSktPr2[0], "mno", 4), 4) will now be called
T41: myRead(daSktPr2[1], Ba, 20) result was 4 Ba: mno
T41: myRead(daSktPr2[1], Ba, 20) will now be called
T32: PE_NOT(myWrite(daSktPr2[0], "mno", 4), 4) result was 4
T32: myClose(daSktPr2[0]) will now be called
T41: myRead(daSktPr2[1], Ba, 20) result was 0
T41: myClose(daSktPr2[1]) will now be called
T41: myClose(daSktPr2[1]) result was 0
T41: myClose(daSktPr1[1]) will now be called
T41: myClose(daSktPr1[1]) result was 0
T41: myClose(daSktPr[1]) will now be called
T41: myClose(daSktPr[1]) result was 0
T41: myClose(daSktPr[1]) will now be called
T41: myClose(daSktPr[1]) result was -1 errno 9: Bad file descriptor
T41: myRead(daSktPr[1], Ba, 20) will now be called
T41: myRead(daSktPr[1], Ba, 20) result was -1 errno 9: Bad file descriptor
T41: myReadcond(daSktPr[1], Ba, 20, 0, 0, 0) will now be called
T41: myReadcond(daSktPr[1], Ba, 20, 0, 0, 0) result was -1 errno 9: Bad file
descriptor
T41: myWrite(daSktPr[1], Ba, 20) will now be called
T41: myWrite(daSktPr[1], Ba, 20) result was -1 errno 9: Bad file descriptor
T41: myTcdrain(daSktPr[1]) will now be called
T41: myTcdrain(daSktPr[1]) result was -1 errno 9: Bad file descriptor
T42: threadT41.join(); statement has now finished
T32: myClose(daSktPr2[0]) result was 0
T32: threadT42.join(); statement will now be started
T32: threadT42.join(); statement has now finished
Pri: T32.join(); statement has now finished
```

- write socket pair using circular buffer

Operation of lock-free circular buffer



- you don't need mutex
 - you can write and read at the same time .
 - interrupt service routine is not allowed to lock the mutex but thread is
- of 2 re-position pointer
(2) and (B)
- ④ if you do (B) re-position pointer first, the reader may misunderstand and keep reading, up to "back" block, even up to "write data" block and end at (B): re-position pointer

351 Sample Midterm Exam

Enter “D100” on the computer answer sheet under “SECTION” and “0001” under “SPECIAL CODE”. When you are finished with the exam, hand in these exam pages, the dataServer.c code, as well as the computer sheet.

Please refer to the distributed code “dataServer.c” for questions A, B, and 01 through 10. This code is **based** on the time1.c code from the Krten textbook, but has significant differences. There are two written questions (marks as specified) and 12 multiple-choice questions (1 mark each), for a total of 20 marks.

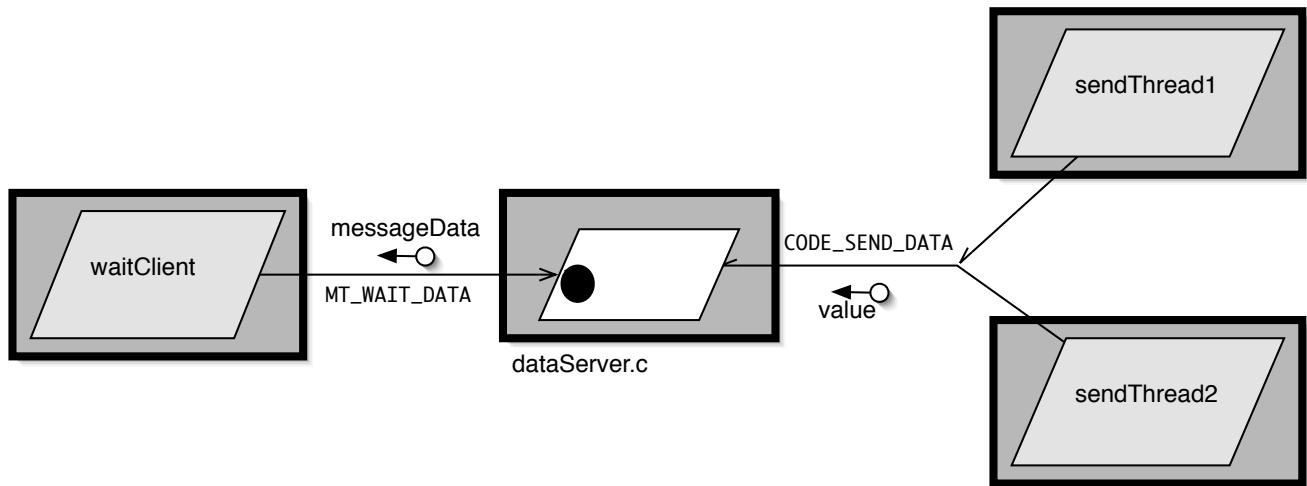
A. Consider a single-threaded client that has a loop in which it will request data from the server, and when it

```
for (;;) {
    clientMessageT msg, rmsg;
    msg.messageType = MT_WAIT_DATA;
    MsgSend(coid, &msg, sizeof(msg), &rmsg, sizeof(rmsg));
    // error handling could go here if MsgSend returned -1
    // if (rmsg.messageType == MT_OK)
        cout << rmsg.messageData << endl;
}
```

loop
3 marks

B. Draw using the collaboration graph notation the server, one data-requesting client, and two data-supplying threads (you may assume each is in a separate process if you wish). Indicate send blocking or reply blocking if either is to be expected. (5 marks)

CGN diagram for the midterm.



Answer each of the following multiple-choice questions. Circle your choice for each question on these pages and eventually fill in the computer answer sheets.

01. In function *main()*, the line “exit (EXIT_FAILURE)” does what?

- a. Ends the primary thread and then immediately ends the process.**
- b. Ends the process but the primary thread still remains.
- c. Ends the primary thread but the process still remains.
- d. Calls the *exit()* function after which the program goes on to the *receiveMessages()* function.
- e. Exits the block of code associated with the “if” statement.

For questions 2 through 10... In function *main()*, consider if the line

`pthread_create (NULL, NULL, receiveMessages, NULL);`

was put before the line

`receiveMessages (NULL);`

(i.e. “uncomment” the *pthread_create()* line)

02. How many threads would now be in the process:

- a. 0
- b. 1
- c. 2**
- d. 3
- e. more than 3

03. Which variable and/or data structures would **have** to be protected from concurrent access to ensure correct operation of the program:

- a. msg
- b. clients[0], clients[1], ..., clients[MAX_CLIENT - 1]**
- c. chid
- d. both b and c
- e. none of the above

04. Of the following choices, with QNX what would be the most-efficient way to protect any variables or data structures needing protection:

- a. use one or more mutexes**
- b. use one or more semaphores
- c. use one or more barriers
- d. use message passing
- e. No variables or data structures need to be protected

05. If a semaphore named `sem4` and initialized to 1 was being used in the program to protect the consistency of one or more variables and/or data structures, what modified code would allow correct operation of these lines from function `gotAMessage()`:

```
for (i = 0; i < MAX_CLIENT; i++) {

    if (!clients [i].in_use) {

        // found one -- mark as in use, save rcvid
        clients [i].in_use = 1;
        clients [i].rcvid = rcvid;
        return;
    }
}
```

a.

```
sem_wait(&sem4);
for (i = 0; i < MAX_CLIENT; i++) {
    if (!clients [i].in_use) {
        clients [i].in_use = 1;
        clients [i].rcvid = rcvid;
        return;
    }
}
sem_post(&sem4);
```

b.

```
for (i = 0; i < MAX_CLIENT; i++) {
    sem_post(&sem4);
    if (!clients [i].in_use) {
        clients [i].in_use = 1;
        clients [i].rcvid = rcvid;
        return;
    }
    sem_wait(&sem4);
}
```

c.

```

for (i = 0; i < MAX_CLIENT; i++) {
    sem_post(&sem4);
    if (!clients [i].in_use) {
        clients [i].in_use = 1;
        clients [i].rcvid = rcvid;
        sem_wait(&sem4);
        return;
    }
    sem_wait(&sem4);
}

```

d.

```

for (i = 0; i < MAX_CLIENT; i++) {
    sem_wait(&sem4);
    if (!clients [i].in_use) {
        clients [i].in_use = 1;
        clients [i].rcvid = rcvid;
        sem_post(&sem4);
        return;
    }
    sem_post(&sem4);
}

```

e.

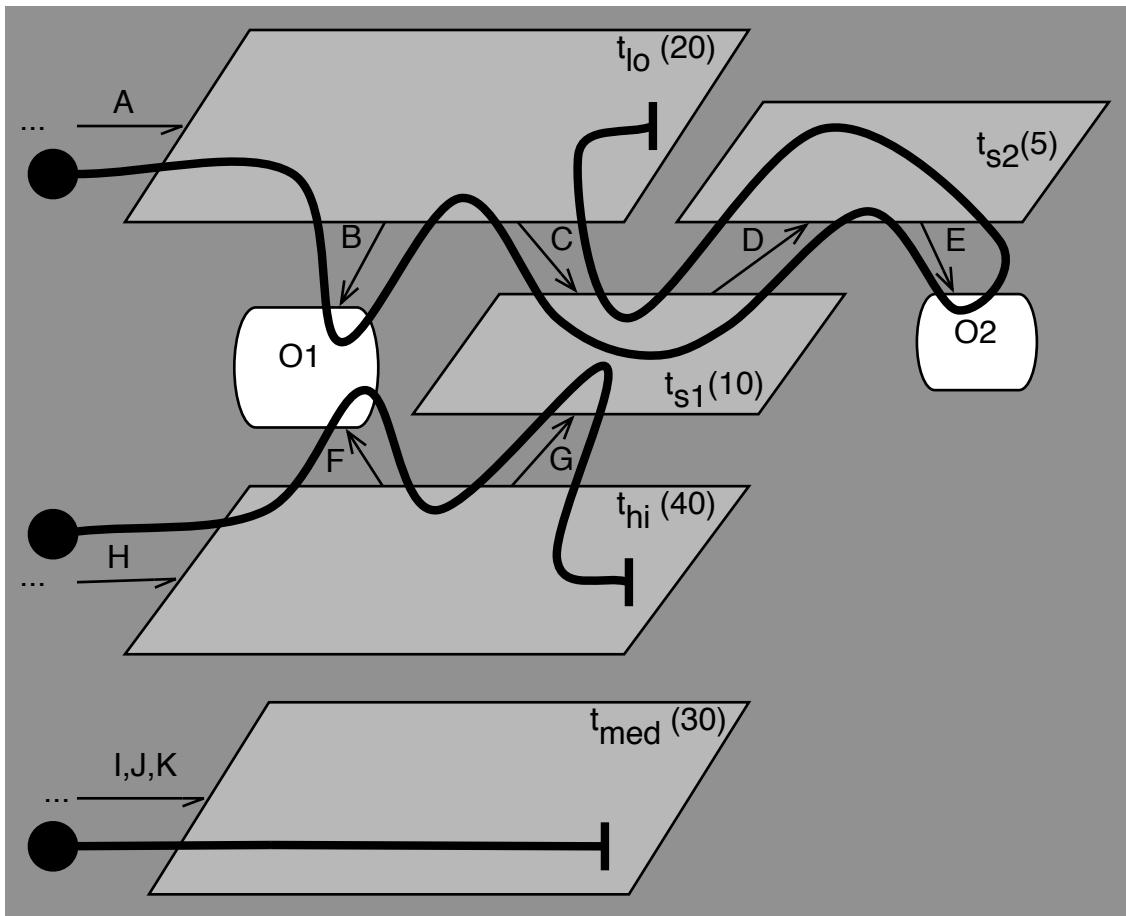
```

for (i = 0; i < MAX_CLIENT; i++) {
    sem_wait(&sem4);
    if (!clients [i].in_use) {
        sem_wait(&sem4);
        clients [i].in_use = 1;
        clients [i].rcvid = rcvid;
        sem_post(&sem4);
        return;
    }
    sem_post(&sem4);
}

```

06. Comparing the program resulting from a correct modification from question 5 (and any other needed modifications) with the original program (i.e. the program exactly as was given out with this exam), on average how much time would transpire using the modified program between receiving a synchronous (i.e. non-pulse) message and replying to it? Assume the programs are being run on a **uniprocessor** system.
- a. roughly half the time it would take the original program
 - b. a little less time than it would take the original program
 - c. exactly the same time it would take the original program
 - d. a little more time than it would take original program**
 - e. roughly twice the time it would take the original program
07. Using for both programs a **very lightly loaded 8-processor symmetric multiprocessor (SMP)** system, on average how much time would transpire using the modified program between receiving a synchronous message and replying to it?
- a. roughly 1/8 of the time it would take the original program
 - b. roughly half the time it would take the original program
 - c. roughly the same time it would take the original program**
 - d. roughly twice the time it would take the original program
 - e. roughly 8 times the time it would take the original program
08. Considering FIFO and standard Round Robin scheduling for the resulting program ...
- a. If the threads in the program use standard Round Robin scheduling, more context switching would typically result.
 - b. The choice of FIFO or standard Round Robin would typically make no difference to this program.**
 - c. Round Robin scheduling would be the best choice for this program.
 - d. both a and c are correct
 - e. none of the above are correct
09. If we turn off priority inheritance and if the thread created by the `pthread_create()` call is given a lower priority than the primary thread ...
- a. on a uniprocessor system the created thread would never be given the chance to start running
 - b. on a lightly loaded symmetric multiprocessor system the created thread would never be given the chance to start running
 - c. the created thread would receive about half the pulses and messages on a lightly loaded multiprocessor
 - d. both a and b are correct
 - e. none of the above are correct**
10. Considering `pthread_join()`, `pthread_barrier_init()`, and `pthread_barrier_wait()` ...
- a. One or more `pthread_join()` calls are needed for this program to run properly.
 - b. One or more `pthread_barrier_init()` calls are needed for this program to run properly.
 - c. One or more `pthread_barrier_wait()` calls are needed for this program to run properly.
 - d. both b and c are correct.
 - e. none of the above are correct**

Refer to the following diagram for the remaining questions.



11. Assuming default QNX priority inheritance behaviour and that the priority of threads when created are shown in parentheses, what priority will thread t_{s1} be running at immediately after the sequence A,B,C,I,H,F

a. 10.

b. 20.

c. 30.

d. 40.

e. None of the above.

12. With the previous assumptions, what priority will thread t_{s1} be running at immediately after the sequence A,B,C,I,H,F,G

a. 10.

b. 20.

c. 30.

d. 40.

e. None of the above.